

Task 1.1. Supervised Learning: Standard Classifier

Garoe Dorta-Perez
CM50246: Machine Learning and AI

November 12, 2014

1 Introduction

Given pictures from the world and been asked to classify them in several groups, we are faced with a problem of multi-class classification. One of the options would be to create N one-against-all binary classifiers. Using x to denote our data, ω to for the world state, and λ for the probability of observing the given class. However a better one involves using a categorical distribution to model our world. Where $\boldsymbol{\lambda}$ is a vector that contains a λ for each class.

$$Pr(\omega|\mathbf{x}) = Cat_w[\boldsymbol{\lambda}[\mathbf{x}]] \quad (1)$$

2 The problem

As stated in the introduction, we are going to fit a Categorical probability model into our data. Denoting I as the total number of data points that we are given. Then, using Bayes' rule we have:

$$Pr(\theta|x_{1...I}) = \frac{\prod_{i=1}^I Pr(\omega = k_n|x, \theta)Pr(\theta)}{Pr(x_{1...I})} \quad (2)$$

We need N activation functions to enforce the constrains. Since we are solving for multi-class classification a logistic sigmoid function as activation will not be valid. Therefore a softmax function is used instead for each activation a_n .

$$a_n = \phi_n^T x \quad (3)$$

$$\lambda_n = softmax_n[a_1, a_2 \cdots a_N] = \frac{exp[a_n]}{\sum_{m=1}^N exp[a_m]} \quad (4)$$

For the Prior we are going to use a Normal distribution with zero mean and σ variance. In order to simplify the calculations we are going to minimise the log of the probability, where y_{in} is the softmax expression for class n and data i :

$$L = -\log \sum_{i=1}^I y_{in} + \frac{1}{2\sigma^2} \phi^T \phi \quad (5)$$

With gradient and Hessian updates being:

$$\begin{aligned} \frac{\delta L}{\delta \phi_n} &= \sum_{i=1}^I (y_{in} - \delta [\omega_i - n]) \mathbf{x}_i + \frac{\phi}{\sigma^2} \\ \frac{\delta^2 L}{\delta \phi_m \delta \phi_n} &= \sum_{i=1}^I y_{im} (\delta [m - n] - y_{in} \mathbf{x}_i \mathbf{x}_i^T) + \frac{\delta [m - n]}{\sigma^2} \end{aligned} \quad (6)$$

To make the predictions we evaluate a new sample doing a Laplace approximation and then a Monte Carlo integration

$$predictions = \int y_{in} \mathcal{N}_a(\mu_a, \Sigma_a) da \quad (7)$$

3 Results

Digits							
Prior Variance	1	10	100	1000	1	1	1
Initial phi	0.1	0.1	0.1	0.1	-1	1	2.5
Accuray	88%	86%	77%	48%	87%	87%	83%

ETH-80-HoG							
Prior Variance	1	10	100	1000	1000	1000	1000
Initial phi	0.125	0.125	0.125	0.125	-10	-1	10
Accuray	67%	74%	84%	89%	89%	89%	89%

Digits initial phi 2.5 works, 5 or bigger runs on NaN error, -10 gives NaN too, -1 works Digits dataset, with prior 100, initial phi ones Elapsed time is 414.619122 seconds. Hits: 71.87

4 Conclusion