

Task 1.2. Supervised Learning: Bayesian Linear Regressor

Garoe Dorta-Perez
CM50246: Machine Learning and AI

December 10, 2014

1 Introduction

SVM are a popular choice in classification problems. However they are inherently not probabilistic. A technique that uses the same intuition is the Relevance Vector Machine. It uses a full Bayesian approach with a prior that encourages sparseness.

2 The problem

First a dual model is used, previously we would model the data as a normal distribution. Where w is the world state, x is the data point, ϕ are the parameters of a linear function of the data and σ is the standard deviation. We now apply a transformation on the parameters as shown in Equation 1, so ϕ is now a weighted sum over the observed data points, where ψ is a vector with the weights. If we are working with few data of high dimensionality, this sum is faster to calculate, as it has less terms, than the original linear dependency.

$$Pr(w_i|\mathbf{x}_i) = Norm_{x_i}[\phi^T \mathbf{x}_i, \sigma^2], \quad (1)$$

$$\phi = \mathbf{X}\psi, \quad (2)$$

Using the dual parameters ψ , we encourage sparseness in the model by using the prior defined in Equation 3. Where I is the number of data points, $Stud$ is a Student's t-distribution with ν degrees of freedom and Γ is the Gamma function. Using a product of t-distributions produces the desired sparseness since the areas with higher probability density are in the origin and along the axis.

$$Pr(\psi) = \prod_{i=1}^I Stud_{\psi_i}[0, 1, \nu], \quad (3)$$

$$= \prod_{i=1}^I \frac{\Gamma(\frac{\nu+1}{2})}{\sqrt{\nu\pi}\Gamma(\frac{\nu}{2})} \left(1 + \frac{\psi_i^2}{\nu}\right)^{-(\nu+1)/2}, \quad (4)$$

A t-distribution is not conjugate to a Normal distribution, so there is no simple closed form solution for the posterior. The solution will be to approximate the t-distributions by maximizing with respect to their hidden variables h , introduced in Equation 5, where \mathbf{w} is a one dimensional array with the world state. This leads to the marginal likelihood shown in Equation 6, where \mathbf{I} is the identity matrix, \mathbf{X} is a matrix with the data points and \mathbf{H} is a matrix with all the hidden variables.

$$Pr(\psi) = \prod_{i=1}^I \int Norm_{\mathbf{w}}[0, 1/h_i] Gam_{h_i}[\nu/2, \nu/2] dh_i, \quad (5)$$

$$Pr(\mathbf{w}|\mathbf{X}, \sigma^2) \approx \max_{\mathbf{H}} \left[Norm_{\mathbf{w}} \left[\mathbf{0}, \mathbf{X}^T \mathbf{X} \mathbf{H}^{-1} \mathbf{X}^T \mathbf{X} + \sigma^2 \mathbf{I} \right] \prod_{d=1}^D Gam_{h_d} [\nu/2, \nu/2] \right], \quad (6)$$

The optimization is performed in three steps:

1. Optimize the marginal likelihood with respect to the hidden variables, using Equation 7.

$$h_i^{new} = \frac{1 - h_i \sum_{ii} + \nu}{\mu_i^2 + \nu}, \quad (7)$$

2. Update $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$, using Equation 8.

$$\begin{aligned} \boldsymbol{\mu} &= \frac{1}{\sigma^2} \mathbf{A}^{-1} \mathbf{X}^T \mathbf{X} \mathbf{w}, \\ \boldsymbol{\Sigma} &= \mathbf{A}^{-1}, \\ \mathbf{A} &= \frac{1}{\sigma^2} \mathbf{X}^T \mathbf{X} \mathbf{X}^T \mathbf{X} + \mathbf{H}, \end{aligned} \quad (8)$$

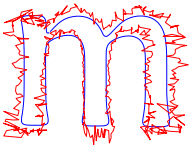
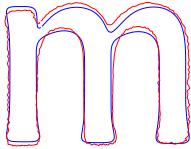
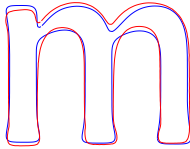
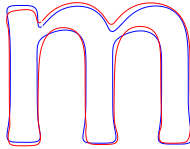
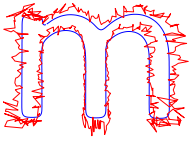
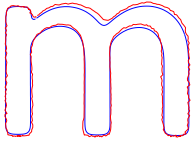
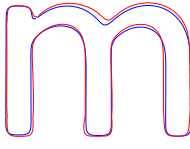
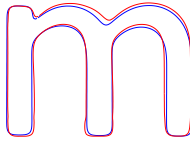
3. Optimize the marginal likelihood with respect to the variance parameter, using Equation 9.

$$(\sigma^2)^{new} = \frac{1}{\mathbf{I} - \sum_i (1 - h_i \sum_{ii})} (\mathbf{w} - \mathbf{X}^T \mathbf{X} \boldsymbol{\mu})^T (\mathbf{w} - \mathbf{X}^T \mathbf{X} \boldsymbol{\mu}), \quad (9)$$

After the training process, we only take the data points whose hidden variable h_i is smaller than a threshold. Since a larger one means a small ψ_i , and hence no significance in the weighted sum. In the implementation, we use a Gaussian kernel $k[x_i, x_j]$, as shown in Equation 10, instead of calculating the inner products $\mathbf{X}^T \mathbf{X}$, where λ controls the scale of the output. This is more efficient than computing the inner product.

$$k[x_i, x_j] = \exp \left[-\frac{1}{2} \left(\frac{(x_i - x_j)^T (x_i - x_j)}{\lambda^2} \right) \right], \quad (10)$$

3 Results

	$\sigma^2 = 10$	$\sigma^2 = 1$	$\sigma^2 = 0.01$	$\sigma^2 = 0.001$
Font1				
Font2				
Font3	